

EECS598-001
Approximation Algorithms & Hardness of Approximation

Pingbang Hu

September 13, 2022

Abstract

This is an advanced graduate-level algorithm course taught in University of Michigan by [Euiwoong Lee](#). Topics include both approximation algorithms like covering, clustering, network design, and constraint satisfaction problems (the first half), and also the hardness of approximation algorithms (the second half).

The first half of the course is classical and well-studied, and we'll use Williamson and Shmoys[[WS11](#)], Vazirani[[Vaz02](#)] as our reference. The second half of the course is still developing, and we'll look into papers by Barak and Steurer[[BS14](#)], O'donnell[[ODo21](#)], etc.

This course is taken in Fall 2022, and the date on the covering page is the last updated time.

Contents

1	Introduction	2
1.1	Computational Problem	2
1.2	Efficient Algorithms	2
1.3	Approximation Algorithms	3
1.4	Hardness	4
2	Set Cover	5
2.1	Covering	5
2.2	Greedy Method	6
2.3	Linear Programming Rounding	7
2.4	Covering-Packing Duality	10
2.5	Primal-Dual Method	11
2.6	Feedback Vertex Set	12

Chapter 1

Introduction

Lecture 1: Overview, Set Cover

1.1 Computational Problem

29 Aug. 10:30

We're interested in the following optimization problem: Given a problem with an input, we want to either maximize or minimize some objectives. This suggests the following definition.

Definition 1.1.1 (Computational problem). A *computational problem* P is a function from input I to (X, f) , where X is the feasible set of I and f is the objective function.

We see that by replacing f with $-f$, we can unify the notion and only consider either minimization or maximization, but we will not bother to do this.

Example (s - t shortest path). The s - t shortest path problem P can be formalized as follows. Given input I , it defines

- Input: Graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ and two vertices $s, t \in \mathcal{V}$.
- Feasible set: $X = \{\text{set of all (simple) paths } s \text{ to } t\}$.
- Objective function: $f: X \rightarrow \mathbb{R}$ where $f(x) = \text{length}(x)$ (# of edges of x).

The output of P should be some $x \in X$ (i.e., some valid s - t paths) such that it minimizes $f(x)$.

We see that the [computational problem](#) we focus on is an optimization problem, and more specifically, we're interested in [combinatorial optimization](#).

Definition 1.1.2 (Combinatorial optimization). A *combinatorial optimization* problem is a [problem](#) where the feasible set X is a finite set.

Example (s - t shortest path). The s - t shortest path problem is an [combinatorial optimization](#) problem since given a graph \mathcal{G} with $n = |\mathcal{V}|$, $m = |\mathcal{E}|$, there are at most $n!$ different paths, i.e., $|X| \leq n! < \infty$.

Note. We'll also look into some continuous optimization problem, where X is now infinite (or even uncountable). For example, find $x \in \mathbb{R}$ that minimizes $f(x) = x^2 + 2x + 1$. In this case, $X = \mathbb{R}$ which is uncountable (hence infinite).

1.2 Efficient Algorithms

Given a [problem](#) P , we want to solve it fast with [algorithms](#). Before we characterize the speed of an [algorithm](#), we should first define what exactly an algorithm is.

Definition 1.2.1 (Algorithm). Given a **problem** P and input I (which defines X and f), an *algorithm* A outputs solution $y = A(I)$ such that $y \in X$ and $y = \arg \max_{x \in X} f(x)$ or $\arg \min_{x \in X} f(x)$, depending on I .

Definition 1.2.2 (Efficient). We say that an **algorithm** A is *efficient* if it runs in **polynomial time**.

Remark (Runtime parametrization). The *runtime* of an **algorithm** A should be parametrized by the size of input I . Formally, given input I represented in s bits, runtime of A on I should be $\text{poly}(s)$ for A to be **efficient**.

Note. In most cases, there are 1 or 2 parameters that essentially define the size of input.

Example (Graph). A natural representation of a graph with n vertices and m edges are

- (a) Adjacency matrix: n^2 numbers.
- (b) Adjacency list: $O(m + n)$ numbers.

Example (Set system). A **set system** with n elements and m sets has a natural representation which uses $O(nm)$ numbers.

Example. If an input I can be represented by s bits, then the runtime of an **algorithm** can be $O(s \log s)$, $O(s^2)$, or $O(s^{100})$, which are considered as **efficient**. On the other hand, something like 2^s or $s!$ are not.

Hence, our goal is to get $\text{poly}((n, m))$ -time **algorithm**!

1.3 Approximation Algorithms

We first note that many interesting **combinatorial optimization problems** are NP-hard, hence it's impossible to find optimum in polynomial time unless P is NP. This suggests one problem: *How well can we do in polynomial time?*

In normal cases, we may assume that objective function value is always positive, i.e., $f: X \rightarrow \mathbb{R}^+ \cup \{0\}$. Then, we have the following definition which characterizes the *slackness*.

Definition 1.3.1 (Approximation algorithm). Given an **algorithm** A , we say A is an α -*approximation algorithm* for a **problem** P if for every input I of P ,

- Min: $f(A(I)) \leq \alpha \cdot \text{OPT}(I)$ for $\alpha \geq 1$
- Max: $f(A(I)) \geq \alpha \cdot \text{OPT}(I)$ for $\alpha \leq 1$

where we define $\text{OPT}(I)$ as $\max_{x \in X} f(x)$ for maximization, $\min_{x \in X} f(x)$ if minimization.

We see that α characterizes the slackness allowed for our **algorithm** A . Now, we're ready to look at some interesting **problems**. Broadly, there are around 10 classes of them which are actively studied:

- We'll see cover, clustering, network design, and constraint satisfaction problems.
- We'll not see: graph cuts, Packing, Scheduling, String, etc.

The above list is growing! For example, applications of continuous optimization in **combinatorial optimization** is getting attention recently. Also, there are around 8 techniques developed, e.g., greedy, local search, LP rounding, SDP rounding, primal-dual, cuts and metrics, etc.

1.4 Hardness

For most problems we saw, we can even say that getting an α -approximation is NP-hard for some $\alpha > 1$. This bound is sometimes tight, but not always, and we'll focus on this part in the second half of this course.

Chapter 2

Set Cover

2.1 Covering

Before we jump into any problem formulations, we define a fundamental object in combinatorial optimization, the [set system](#).

Definition 2.1.1 (Set system). Given a ground set Ω (often called *universe*), the *set system* is an order pair (Ω, \mathcal{S}) where \mathcal{S} is a collection of subsets of Ω .

Note. For a [set system](#) (Ω, \mathcal{S}) , we often let $m := |\mathcal{S}|$ and $n := |\Omega|$.

Definition 2.1.2 (Degree). Given a [set system](#) (Ω, \mathcal{S}) , the *degree* of $x \in \Omega$, $\deg(x)$, is defined as

$$\deg(x) := |\{S \in \mathcal{S} \mid x \in S\}|.$$

Remark (Bipartite representation). Naturally, for a [set system](#), we have a bipartite representation.




Figure 2.1: Bipartite representation of a [set system](#).

Denote $d := \max_{e \in U} \deg(e) \leq m$ and $k := \max_{i \in [m]} |S_i| \leq n$, which is just the maximum vertex degree on two sides of the bipartite graph representation of this [set system](#).

Finally, we have the following.

Definition 2.1.3 (Covering). A *covering* $\mathcal{S}' \subseteq \mathcal{S}$ of (Ω, \mathcal{S}) is a (sub)collection of subsets such that $\bigcup_{S \in \mathcal{S}'} S = \Omega$.

Let's first consider the classical problem called [set cover](#).

Problem 2.1.1 (Set cover). Given a finite [set system](#) (U, \mathcal{S}) where $\mathcal{S} := \{S_i \subseteq U\}_{i=1}^m$ along with a weight function $w: \mathcal{S} \rightarrow \mathbb{R}^+$, find a [covering](#) \mathcal{S}' while minimizing $\sum_{S \in \mathcal{S}'} w(S)$.

Assuming there always exists at least one [covering](#), we can in fact get two types of non-comparable approximation ratio in terms of k and d . Specifically, we get $\log k$ and d -approximation ratio via either greedy, LP rounding or dual-methods.

2.2 Greedy Method

We first see the algorithm when $w(S) = 1$ for all $S \in \mathcal{S}$.

Algorithm 1: [Set cover](#) – Greedy

Data: A [set system](#) (U, \mathcal{S})

Result: A [covering](#) \mathcal{S}'

```

1  $\mathcal{S}' \leftarrow \emptyset, i \leftarrow 0$ 
2 while  $U \neq \emptyset$  do                                     //  $O(n)$ 
3   Choose  $S_i$  with maximum  $|U \cap S_i|$                    //  $O(mn)$ 
4    $\mathcal{S}' \leftarrow \mathcal{S}' \cup \{S_i\}$ 
5    $U \leftarrow U \setminus S_i$ 
6   for  $e \in U \cap S_i$  do
7      $y_e \leftarrow w(S_i) / |U \cap S_i|$                  // Average costs
8    $i \leftarrow i + 1$ 
9 return  $\mathcal{S}'$ 

```

We focus on the case that $w(S) = 1$ for all S .

Remark. It's clear that [Algorithm 1](#) is a polynomial time algorithm, also, the output \mathcal{S}' is always a valid [covering](#).

Theorem 2.2.1. [Algorithm 1](#) is an H_k -approximation^a algorithm.

^a H_k is the so-called *harmonic number*, which is defined as $\sum_{i=1}^k 1/i \leq \ln k + 1$.

Proof. Denote the OPT as $\mathcal{S}^* := \{S_1^*, \dots, S_\ell^*\}$, and first note that the average cost y_e essentially maintains $\sum_{e \in U} y_e = |\mathcal{S}'|$, hence we just need to bound y_e w.r.t. \mathcal{S}^* . To do this, for any $S^* \in \mathcal{S}^*$, say $S_1^* = \{e_1, \dots, e_k\}$ where we number e_i in terms of the order of which being deleted, i.e., e_1 is deleted first from U ([line 5](#)), etc.

Note. S_1^* can have less than k element, but in that case similar argument will follow. Also, if some elements are deleted at the same time, we just order them arbitrarily.

Then, we have the following claim.

Claim. For all e_i , $y_{e_i} \leq \frac{1}{k-i+1}$.

Proof. Consider the iteration when e_i was picked by \mathcal{S}' , i.e., $|U \cap \mathcal{S}'| \geq |U \cap S_1^*| \geq k - i + 1$, then by definition ([line 5](#)) we have $y_{e_i} = \frac{1}{|U \cap \mathcal{S}'|} \leq \frac{1}{|U \cap S_1^*|} \leq \frac{1}{k-i+1}$. ⊗

We immediately see that whenever the optimal solution pays 1 (for choosing S_1^* for instance), [Algorithm 1](#) pays at most H_k since $\sum_{e_i \in S_1^*} y_{e_i} \leq \sum_{i=1}^k \frac{1}{k-i+1} = H_k$, or more formally,

$$|\mathcal{S}'| = \sum_{e \in U} y_e \leq \sum_{S_i^* \in \mathcal{S}^*} \underbrace{\sum_{e \in S_i^*} y_e}_{\leq H_k} \leq \ell \cdot H_k = H_k \cdot |\text{OPT}|,$$

which finishes the proof. ■

In all, observe that $H_k \leq \ln k + 1$, we see that [Algorithm 1](#) is a $(\ln k)$ -approximation algorithm. Also, the weighted version can be easily derived by replacing 1 with the corresponding weight.

Lecture 2: Linear Programming with Set Covers

2.3 Linear Programming Rounding

31 Aug. 10:30

To get a d -approximation algorithm, instead of seeing the greedy algorithm, we first see the LP¹ dual method, which turns out to be exactly the same as the greedy algorithm.

As previously seen. Both linear programming and convex programming can be solved in polynomial time.

Notice that it's more natural to define **set cover** in terms of ILP (integer LP). Define our integer variables $\{x_i\}_{i \in [n]}$ such that

$$x_i = \begin{cases} 1, & \text{if } S_i \in \mathcal{S}'; \\ 0, & \text{otherwise.} \end{cases}$$

In this way, we have the following ILP formulation for **set cover** as

$$\begin{aligned} \min \quad & \sum_i w_i \cdot x_i \\ & \sum_{S_i \ni e} x_i \geq 1 & \forall e \in U \\ \text{(IP)} \quad & x_i \in \{0, 1\} & \forall i. \end{aligned}$$

But we know that this is an NP-hard problem, so we relax it to be

$$\begin{aligned} \min \quad & \sum_i w_i \cdot x_i \\ & \sum_{S_i \ni e} x_i \geq 1 & \forall e \in U \\ \text{(LP)} \quad & x_i \geq 0 & \forall i. \end{aligned}$$

Write it in a more compact form, we have

$$\begin{aligned} \min \quad & \langle w, x \rangle \\ & Ax \geq \mathbb{1} \\ & x \geq 0 \end{aligned}$$

where $A \in \mathbb{R}^{n \times m}$ such that

$$A_{ij} = \begin{cases} 1, & \text{if } e_i \in S_j; \\ 0, & \text{otherwise.} \end{cases}$$

Note. Note when we do relaxation, we want $x \in \text{fes}(\text{IP}) \Rightarrow x \in \text{fea}(\text{LP})$, i.e., $\text{fes}(\text{LP}) \supseteq \text{fes}(\text{IP})$. Note that in this case, for a minimization problem, we have

$$f(x) = \text{OPT}_{\text{LP}} \leq \text{OPT}_{\text{IP}}.$$

In this case, we see that the most natural way to get an integer solution from the fractional solution obtained from the relaxed LP is to **round** x to integral solution. This leads to the following **algorithm**.

Algorithm 2: **Set cover** – LP Rounding

Data: A **set system** (U, \mathcal{S})

Result: A **covering** \mathcal{S}'

```

1  $x \leftarrow \text{solve}(\text{LP})$  // Solve the relaxed (LP)
2  $\mathcal{S}' \leftarrow \{S_i : x_i \geq 1/d\}$ 
3 return  $\mathcal{S}'$ 
```

We now prove the correctness and **Algorithm 2**'s approximation ratio.

¹See **MATH561** for a complete reference.

Lemma 2.3.1. \mathcal{S}' is a covering.

Proof. Fix $e \in U$, let S_1, \dots, S_d be the sets containing e . We see that

$$\sum_{i=1}^d x_i \geq 1 \Rightarrow \exists j \in [d] \text{ s.t. } x_j \geq \frac{1}{d} \Rightarrow S_j \in \mathcal{S}'.$$

■

Theorem 2.3.1. Algorithm 2 is d -approximation algorithm.

Proof. By comparing $w(\mathcal{S}')$ and $\text{OPT}_{\text{LP}} = \sum_{i=1}^m x_i w_i$, we see that

$$\text{OPT} \leq \sum_{S_i \in \mathcal{S}'} w_i \leq d \sum_{S_i \in \mathcal{S}'} w_i x_i \leq d \cdot \text{OPT}_{\text{LP}} \leq d \cdot \text{OPT},$$

which implies $\text{OPT} / d \leq \text{OPT}_{\text{LP}} \leq \text{OPT}$.

Note. Note that OPT is assumed to be OPT_{IP} , i.e., the optimum of the original IP formulation of Problem 2.1.1.

■

Definition 2.3.1 (Integrality gap). The *integrality gap* as

$$\sup_{\text{input } I} \frac{\text{OPT}(I)}{\text{OPT}_{\text{LP}}(I)}.$$

Remark. We see that the *integrality gap* of Algorithm 2 is d from Theorem 2.3.1.

2.3.1 Randomized Linear Programming Rounding

And indeed, we can use a more natural way to do the rounding, i.e., respect to the x_i value.

Intuition. If x_i is close to 1, it's reasonable to include it, vice versa.

We see that algorithm first.

Algorithm 3: Set cover – Randomized LP Rounding

Data: A set system (U, \mathcal{S})

Result: A (possible) covering \mathcal{S}'

```

1  $x \leftarrow \text{solve}(\text{LP})$  // Solve the relaxed (LP)
2  $\mathcal{S} \leftarrow \emptyset$ 
3 for  $i = 1, \dots, m$  do
4    $\lfloor$  add  $S_i$  to  $\mathcal{S}'$  w.p.  $x_i$  // independently
5 return  $\mathcal{S}'$ 
```

Now, the question is, how is this \mathcal{S}' 's quality? In other words, fix $e \in U$, what's $\Pr(e \text{ is covered})$?

Lemma 2.3.2. $\Pr(e \text{ is covered}) \geq 1 - 1/e \approx 0.63$.

Proof. We bound $\Pr(\overline{e \text{ is covered}})$ instead. Say S_1, \dots, S_d are the sets containing e , then we see that

$$\Pr(\overline{e \text{ is covered}}) = \prod_{i=1}^d (1 - x_i) \leq \prod_{i=1}^d e^{-x_i} = e^{-\overbrace{(x_1 + \dots + x_d)}^{\geq 1}} \leq e^{-1}.$$

Note. For every x , we have $1 + x \leq e^x$, and this approximation is close when $|x|$ is small. ■

A standard way to boost the correctness of a randomized algorithm is to run it multiple time, which leads to the following.

Algorithm 4: Set cover – Multi-time Randomized LP Rounding

Data: A set system (U, \mathcal{S}) , α
Result: A (possible) covering \mathcal{S}'

```

1  $x \leftarrow \text{solve}(\text{LP})$  // Solve the relaxed (LP)
2  $\mathcal{S} \leftarrow \emptyset$ 
3 for  $t = 1, \dots, \alpha$  do // independently
4   for  $i = 1, \dots, m$  do
5      $\lfloor$  add  $S_i$  to  $\mathcal{S}'$  w.p.  $x_i$  // independently
6 return  $\mathcal{S}'$ 
```

Lemma 2.3.3. With $\alpha = 2 \ln n$, \mathcal{S}' returned from Algorithm 4 is a covering w.p. at least $1 - \frac{1}{n}$.

Proof. We have $\Pr(e \text{ is not covered}) \leq e^{-\alpha}$ from independence of each run. Let $\alpha = 2 \ln n$, then $\Pr(e \text{ is not covered}) \leq e^{-\alpha} = 1/n^2$. By union bound,

$$\Pr(\text{some elements is not covered}) \leq \sum_{e \in U} \Pr(e \text{ not covered}) \leq n \cdot \frac{1}{n^2} = \frac{1}{n}.$$

This implies w.p. at least $1 - 1/n$, \mathcal{S}' is a covering. ■

In other words, with $\alpha = 2 \ln n$, Algorithm 4 is correct with probability at least $1 - 1/n$.

Lemma 2.3.4. With $\alpha = 2 \ln n$, \mathcal{S}' returned from Algorithm 4 has an approximation ratio $4 \ln n$ w.p. at least $\frac{1}{2}$.^a

^aNote that \mathcal{S}' is not necessary a covering.

Proof. Since $\mathbb{E}[w(\mathcal{S}')] \leq \alpha \sum_i w_i x_i = \alpha \text{OPT}_{\text{LP}}$, we have $\Pr(w(\mathcal{S}') \geq 2 \cdot \alpha \text{OPT}_{\text{LP}}) \leq 1/2$ from Markov inequality. We see that w.p. $\geq 1/2$, $w(\mathcal{S}') \leq 2 \cdot 2 \ln n \cdot \text{OPT}_{\text{LP}} \leq 4 \ln n \text{OPT}$. ■

Theorem 2.3.2. By running Algorithm 4 many times, we get a $(4 \ln n)$ -approximation algorithm with high probability.^a

^aNote that we still need to choose \mathcal{S}' .

Proof. Together with Lemma 2.3.3 and Lemma 2.3.4 and using the union bound, the probability of \mathcal{S}' not being a covering or with weight higher than $4 \ln n \text{OPT}$ is at most $\frac{1}{n} + \frac{1}{2}$, which is less than 1. Hence, by running Algorithm 4 many times (independently), the failing possibility is exponential small. ■

Note. With Theorem 2.3.2, we still need to find a valid covering with the lowest cost, where a valid covering with low enough weight is guaranteed to exist with high probability. Note that this is still a polynomial time algorithm since we know that checking \mathcal{S}' is a covering is just linear.

Remark. Indeed, with some smarter algorithm modified from Algorithm 4, we can get a H_k approximation ratio.

Lecture 3: Covering-Packing Duality and Primal-Dual Method

2.4 Covering-Packing Duality

7 Sep. 10:30

We first define some useful notions.

Definition 2.4.1 (Strongly independent). Given a set system (U, \mathcal{S}) , we say $C \subseteq U$ is *strongly independent* if there does not exist $S \in \mathcal{S}$ such that $|C \cap S| \geq 2$.

Remark. Then for any *strongly independent* set $C \subseteq U$, we know that $\text{OPT}_{\text{SC}} \geq |C|$.^a

^aSC denotes *set cover*.

Now, we're trying to find the **strongest witness** of *strongly independent set*, which suggests we define the following problem.

Problem 2.4.1 (Maximum strongly independent set). Given a set system (U, \mathcal{S}) , we want to find the largest *strongly independent set*.

Remark. For any set system, we have $\text{OPT}_{\text{SIS}} \leq \text{OPT}_{\text{SC}}$.^a

^aSIS denotes *maximum strongly independent set*.

As previously seen (LP dual). Recall how we get the dual of a given LP:

$$\begin{array}{ll} \min & c^\top x \\ & Ax \geq b \\ (P) & x \geq 0 \end{array} \qquad \begin{array}{ll} \max & y^\top b \\ & y^\top A \leq c^\top \\ (D) & y \geq 0. \end{array}$$

Also, recall the weak duality ($\text{OPT}_P \geq \text{OPT}_D$) and strong duality ($\text{OPT}_P = \text{OPT}_D$).

Definition 2.4.2 (Covering LP). A primal LP with $A, b, c \geq 0$ is called a *covering LP*.

Definition 2.4.3 (Packing LP). A dual LP with $A, b, c \geq 0$ is called a *packing LP*.

We now give another LP formulation for the unweighted *set cover*. Given $\mathcal{S} = \{S_1, \dots, S_m\}$, $U = \{e_1, \dots, e_n\}$ and define $A \in \mathbb{R}^{n \times m}$ such that

$$A_{ij} = \begin{cases} 1, & \text{if } e_i \in S_j; \\ 0, & \text{otherwise.} \end{cases}$$

Then our LP is defined as

$$\begin{array}{ll} \min & \sum_{j=1}^m x_j \\ & Ax \geq \mathbf{1} \\ (P) & x \geq 0 \end{array} \qquad \begin{array}{ll} \max & \sum_{i=1}^n y_i \\ & y^\top A \leq \mathbf{1} \\ (D) & y \geq 0. \end{array}$$

We see that if we restrict $y_i \in \{0, 1\}$, we see that the dual (D) is just **Problem 2.4.1**. This can be seen via writing the constraint explicitly:

$$\sum_{i=1}^n A_{ij} y_i \leq 1 \Leftrightarrow \sum_{i: e_i \in S_j} y_i \leq 1 \text{ for } j \in [m].$$

And indeed, if we look at the weighted version, we have $\sum_{i: e_i \in S_j} y_i \leq w(S_j)$.

Now, recall the [claim](#) in [Theorem 2.2.1](#), i.e., $y_{e_i} \leq \frac{w(S_j)}{k-i+1}$. We see that the y_{e_i} are just the dual variables in our setup. Additionally, with the observation that we can do this for any set $S = \{e_1, \dots, e_k\} \in \mathcal{S}$, we have the following lemma.

Lemma 2.4.1. The variable $y' := y/H_k$ is dual-feasible, i.e., it's feasible for (D) .

Proof. We see that $y_{e_i} \geq 0$ (and hence y_i) trivially, so we only need to show that

$$\sum_{i=1}^n A_{ij} y' = \sum_{i=1}^n A_{ij} \frac{y_{e_i}}{H_k} \leq w(S_j)$$

for $j \in [m]$. But this is trivial by plugging in $y_{e_i} \leq \frac{w(S_j)}{k-i+1}$ as shown in [Theorem 2.2.1](#), hence

$$\sum_{i=1}^n A_{ij} \frac{y_{e_i}}{H_k} \leq \frac{1}{H_k} \sum_{i=1}^n A_{ij} \frac{w(S_j)}{k-i+1} \leq \frac{1}{H_k} \sum_{i=1}^k \frac{w(S_j)}{k-i+1} = w(S_j),$$

and we're done.^a ■

^aNote that in the above derivation, i is kind of overloading, i.e., e_i corresponding to only some i (confusing, but it's how it is...).

With [Lemma 2.4.1](#), we simply run [Algorithm 1](#) while maintaining y_e for every e , and we're done.

Theorem 2.4.1. [Algorithm 1](#) is an H_k -approximation algorithm in the view of its dual.

Proof. Same as [Theorem 2.2.1](#), but now we have different interpretation. Specifically, if $y' = y/H_k$ is dual-feasible, we know that the corresponding objective value of y' is at most $\text{OPT}_{\text{LP}_D} = \text{OPT}_{\text{LP}_P}$, which is at most OPT_{SC} further. Now, since we're dealing with LP, everything is linear includes the objective value, i.e., y is at most $H_k \cdot \text{OPT}_{\text{SC}}$. ■

Remark (Dual fitting). The above method is called *dual fitting*, which is universal as one can easily see. The way to do this is the following.

1. Given an algorithm, distribute the algorithm to $\{y_i\}$.
2. Prove that y/α is dual-feasible.
3. This shows the algorithm is α -approximation algorithm.

2.5 Primal-Dual Method

We first see the general description of the so-called *primal-dual method*.

1. Maintain x (primal solution) and y (dual solution) where x is integral and infeasible, while y is fractional and feasible. Start from $x = y = 0$.
2. **Somehow** increase y until some dual constraints get tight.
3. **Choose** primal variables correspond to tight dual constraints, and update input accordingly.

Remark. We're using dual variables to get a certificate of the lower bound of the optimal problem we're solving.

In terms of [set cover](#), we have the following.

Algorithm 5: Set cover – Primal-Dual**Data:** A set system (U, \mathcal{S}) **Result:** A covering \mathcal{S}'

```

1  $\mathcal{S}' \leftarrow \emptyset, y \leftarrow 0$ 
2 while  $U \neq \emptyset$  do
3   Choose any  $e \in U$ 
4   Raise  $y_e$  until some constraints get tight
5    $\mathcal{S}' \leftarrow \mathcal{S}' \cup \{\text{sets corresponding to tight dual constraints}\}$ 
6   Update  $U$  // Remove newly covered element in  $U$ 
7 return  $\mathcal{S}'$ 

```

Remark. Algorithm 5 is correct and can be implemented efficiently.

Theorem 2.5.1. Algorithm 5 is a d -approximation algorithm.

Proof. Firstly, y is feasible. And we see that

$$w(\mathcal{S}') = \sum_{S \in \mathcal{S}'} w(S) = \sum_{S \in \mathcal{S}'} \sum_{e \in S} y_e \leq d \cdot \sum_{e \in U} y_e \leq d \cdot \text{OPT}_{\text{LP}_D} = d \cdot \text{OPT}_{\text{LP}_P} \leq d \cdot \text{OPT}_{\text{SC}}.$$

■

Lecture 4: Feedback Vertex Set

2.6 Feedback Vertex Set

12 Sep. 10:30

Following the discussion on primal-dual method, we see another covering problem.

2.6.1 Introduction

We consider the following problem.

Problem 2.6.1 (Feedback vertex set). Given a graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ and a weight function $c: \mathcal{V} \rightarrow \mathbb{R}^+$, we want to find $F \subseteq \mathcal{V}$ with $\min c(F)$ such that $\mathcal{G}[\mathcal{V} \setminus F]$ has no cycle.^a

^aThis is equivalent as saying that $\mathcal{G}[\mathcal{V} \setminus F]$ is a forest.

Note (Feedback). The name *feedback* comes from the fact that if there's a cycle in \mathcal{G} , then it kind of creates *feedback*.

Note (Edge version). The *edge version* of Problem 2.6.1 can be solved by finding $T \subseteq \mathcal{E}$ be the maximum weight forest,^a and let $F := \mathcal{E} \setminus T$.

^aThis can be found exactly in polynomial time.

Notation. In this lecture, when talking about cycle, we're referring to the vertices in which. But the meaning can vary from context to context.

Remark. This is a *special case* of Problem 2.1.1.

Proof. Let $\mathcal{C} := \{\text{set of all (simple) cycles}\}$ and consider Problem 2.1.1 on the set system $(\mathcal{C}, \mathcal{V})$, i.e., we want to find $F \subseteq \mathcal{V}$ such that $\forall C \in \mathcal{C}, |F \cap C| \geq 1$. ⊗

Note. The naive algorithm by directly applying methods discussed for [Problem 2.1.1](#), we see that since $\min(\log k, d) = \Omega(n)$ for k being the maximum set size (which is $2^{\Omega(n)}$) and $d = n$, the approximation ratio we can get is $\Omega(n)$, which depends on the size of the input.

Now, the goal in this section is to show the following.

Theorem 2.6.1. There exists a 4-approximation algorithm for [Problem 2.6.1](#).

Remark. Actually, there exists a 2-approximation algorithm.

We also have a hardness of [Problem 2.6.1](#).

Theorem 2.6.2. A $(2 - \epsilon)$ -approximation algorithm doesn't exist for all $\epsilon > 0$ assuming some conjecture.

2.6.2 Cycle Covering LP

The most natural LP which models [Problem 2.6.1](#) is the so-called *cycle covering LP*, which can be defined as

$$\begin{aligned} \min \quad & \sum_{v \in \mathcal{V}} c(v)x_v \\ & \sum_{v \in C} x_v \geq 1 \quad \forall \text{ cycle } C \in \mathcal{C} \\ & x \geq 0, \end{aligned}$$

with the variables being $\{x_v\}_{v \in \mathcal{V}}$ such that $x_v = \mathbb{1}_{v \in F}$.

Remark. We see that this cycle covering LP has $2^{\Omega(n)}$ constraints. But we can actually solve this and get an $O(\log n)$ -approximation ratio by smartly rounding the solution. And we can show that this approximation ratio is optimal in terms of this particular LP.

2.6.3 Density LP

A more sophisticated LP is the so-called *density LP*, defined as

$$\begin{aligned} \min \quad & \sum_{v \in \mathcal{V}} c(v)x_v \\ & \sum_{v \in S} x_v(d_v^S - 1) \geq |E(S)| - |S| + 1 \quad \forall S \subseteq \mathcal{V} \\ & x \geq 0 \end{aligned}$$

with the variables being $\{x_v\}_{v \in \mathcal{V}}$.

Notation. The $E(S)$ denotes the edge set in the induced graph $\mathcal{G}[S] = (S, E(S))$, while d_v^S denotes the degree of v in $\mathcal{G}[S]$.

Intuition. The constraint is equivalent as saying that for every induced graph, $\#e \leq \#v - 1$, i.e., we require it to be a forest. Explicitly, $S \subseteq \mathcal{V}$,

$$|E(S)| - \sum_{v \in S} x_v d_v^S \leq |S| - \sum_{v \in S} x_v - 1.$$

Note that in the constraint, the right-hand side is just a lower-bound of $\#e$.

We see that the above LP is not exactly a [covering LP](#) since the coefficients can be negative if a set S is not [irreducible](#).

Definition 2.6.1 (Irreducible). The set $S \subseteq \mathcal{V}$ is *irreducible* if for all $v \in S$, v belongs to some cycles in $G[S]$.

Now, it's clear that by looking at $\mathcal{S} = \{S \subseteq \mathcal{V} \mid S \text{ is irreducible}\}$, we have a *covering LP* defined as

$$\begin{aligned} \min \quad & \sum_{v \in \mathcal{V}} c(v)x_v \\ & \sum_{v \in S} x_v(d_v^S - 1) \geq |E(S)| - |S| + 1 =: b_S \quad \forall S \in \mathcal{S} \\ & x \geq 0. \end{aligned}$$

We first see why this LP models [Problem 2.6.1](#).

Lemma 2.6.1. The integer version of density LP (denote as IP) is equivalent to [Problem 2.6.1](#).

Proof. If x is feasible for [Problem 2.6.1](#), then x is feasible for the IP. On the other hand, if x is feasible for IP, then for every cycle $C \in \mathcal{C}$, x deletes at least 1 vertex from C . ■

2.6.4 Primal-Dual Method

Now we're ready to solve this LP via primal-dual method. Denote the dual variables as $\{y_S\}_{S \in \mathcal{S}}$, then the dual is

$$\begin{aligned} \max \quad & \sum_{S \in \mathcal{S}} y_S b_S \\ & \sum_{S \ni v} (d_v^S - 1)y_S \leq c(v) \quad \forall v \in \mathcal{V} \\ & y \geq 0. \end{aligned}$$

Note. For the density LP and its dual, the constraint is still exponentially many, and no one knows how to solve this. But the power of primal-dual method is that we don't really solve this, rather, we just maintain two sets of solutions for both primal and dual. Moreover, we can maintain the primal solution in integral, while the dual solution in fractional.

We now have the following algorithm.

Algorithm 6: [Feedback vertex set](#) – Primal-Dual

Data: A graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$

Result: A minimal [feedback vertex set](#) F'

```

1   $S \leftarrow \mathcal{V}$ ,  $c' = c$ ,  $y \leftarrow 0$                                 //  $c' \in \mathbb{R}^n$  keeps track of slackness of  $c$ 
2   $S \leftarrow \text{reduce}(S)$                                            // Compute  $\{v \in S : v \text{ belongs to some cycles in } \mathcal{G}[S]\}$ 
3
4  while  $S \neq \emptyset$  do
5       $\alpha \leftarrow \min_{v \in S} c'(v)/(d_v^S - 1)$                 // time for  $y_S$  gets tight by increasing unit weight
6       $y_S \leftarrow \alpha$ 
7       $c'(v) \leftarrow c'(v) - \alpha(d_v^S - 1)$ 
8       $Z \leftarrow \{v \in S : c'(v) = 0\}$ 
9       $F \leftarrow F \cup Z$ ,  $S \leftarrow S \setminus Z$ 
10      $S \leftarrow \text{reduce}(S)$ 
11
12  $F' \leftarrow F$                                                     // Compute a minimal feedback vertex set
13 for  $i = \ell, \dots, 1$  do
14     if  $F' \setminus \{v_i\}$  is a feedback vertex set for  $\mathcal{G}$  then
15          $F' \leftarrow F' \setminus \{v_i\}$ 
16 return  $F'$ 

```

Note. F is a **feedback vertex set** and y is dual-feasible.

We see that we want to compare $\sum_{S \in \mathcal{S}} y_S b_S$ and

$$c(F) = \sum_{v \in F} c(v) = \sum_{v \in F} \sum_{S \ni v} (d_v^S - 1) y_S = \sum_{S \in \mathcal{S}} y_S \sum_{v \in F \cap S} (d_v^S - 1).$$

If we have some ratio of $\sum_{v \in F \cap S} (d_v^S - 1)$ and b_S for every S , we're done.

Lemma 2.6.2. For all $S \in \mathcal{S}$, if F is *minimal* in S ,^a then we have

$$\sum_{v \in S \cap F} (d_v^S - 1) \leq 4 \cdot b_S = 4(|E(S)| - |S| + 1).$$

^ai.e., in $\mathcal{G}[S]$, no $F' \subsetneq F \cap S$ in **feedback vertex set**.

Proof. Let's first see a simple case.

Intuition. If the graph is 3-regular, then we see that the left-hand side is $\leq 2 \cdot |S|$, while the right-hand side is $2 \cdot |S| + 1$.

Let A be the set of degree 2 vertices, and B be the set of vertices with degree larger than 3. Now, consider line segment in the graph. If ℓ is a line segment,

- (a) $|F \cap \ell| \leq 1$
- (b) If F contains one of the endpoints of ℓ , then $|F \cap \ell| = 0$

Since F is minimal, the left-hand side is

$$|A \cap F| + \sum_{v \in B \cap F} (d_v^S - 1) \leq \sum_{v \in B \setminus F} d_v^S / 2 + \sum_{v \in B \cap F} (d_v^S - 1) \leq \sum_{v \in B} (d_v^S - 1),$$

while the right-hand side is

$$|E(S)| - |S| = \sum_{v \in S} (d_v^S / 2 - 1) = \sum_{v \in B} (d_v^S / 2 - 1).$$

It's clear that since $\forall v \in B$, $d_v^S - 1 \leq 4(d_v^S / 2 - 1)$, rearranging this inequality gives the result. ■

We see that to show [Theorem 2.6.1](#), it's enough to check that F' is minimal, then the result follows from [Lemma 2.6.1](#). Hence, we now analyze F' . Clearly, F' is still a **feedback vertex set**. The minimality of F' is guaranteed by the following lemma.

Lemma 2.6.3. F' is minimal in every S_i , where $S_i = S$ when v_i is deleted.

Proof. Suppose this is not the case. Then there exists $v_j \in F'$ such that in $\mathcal{G}[S_i]$, $b_S = \cap S_i \setminus \{v_j\}$ is still a **feedback vertex set**. This means $|F \cap S_j| \setminus \{v_j\}$ is still a **feedback vertex set** in $\mathcal{G}[S_j]$.

Observe that there are no $C \ni v_j$ such that $C \setminus S_j$ is disjoint from F , which further implies $F' \setminus \{v_j\}$ is still a **feedback vertex set** in \mathcal{G} when $i = j$ in ?? . ■

Finally, we see that we can prove [Theorem 2.6.1](#).

Proof of Theorem 2.6.1. By designing an algorithm which runs [Algorithm 6](#) and ?? sequentially, we get a 4-approximation of the density IP guaranteed by [Lemma 2.6.2](#) and [Lemma 2.6.3](#).

From [Lemma 2.6.1](#), [Problem 2.6.1](#) and this density IP is equivalent, proving the theorem. ■

Appendix

Bibliography

- [BS14] Boaz Barak and David Steurer. *Sum-of-squares proofs and the quest toward optimal algorithms*. 2014. DOI: [10.48550/ARXIV.1404.5236](https://arxiv.org/abs/1404.5236). URL: <https://arxiv.org/abs/1404.5236>.
- [ODo21] Ryan O'Donnell. *Analysis of Boolean Functions*. 2021. DOI: [10.48550/ARXIV.2105.10386](https://arxiv.org/abs/2105.10386). URL: <https://arxiv.org/abs/2105.10386>.
- [Vaz02] V.V. Vazirani. *Approximation Algorithms*. Springer Berlin Heidelberg, 2002. ISBN: 9783540653677. URL: <https://books.google.com/books?id=EILqAmzKgYIC>.
- [WS11] D.P. Williamson and D.B. Shmoys. *The Design of Approximation Algorithms*. Cambridge University Press, 2011. ISBN: 9781139498173. URL: https://books.google.com/books?id=Cc%5C_Fdqf3bBgC.